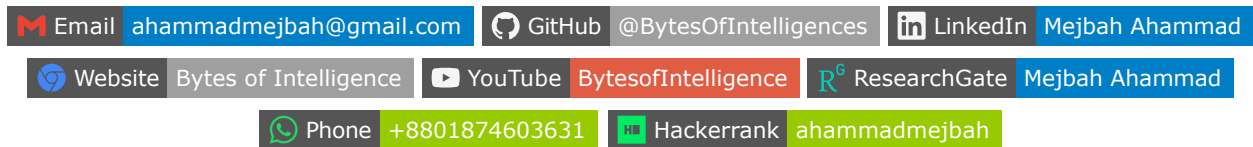




# Introduction to Scikit Learn



## Introduction

Scikit-Learn is a popular open-source machine learning library for Python. It provides simple and efficient tools for data mining and data analysis, accessible to everybody and reusable in various contexts.

## Overview

Scikit-Learn is built upon Python's NumPy and SciPy and is known for its ease-of-use and versatility for various machine learning tasks.

## Key Features

The library includes:

- Simple and efficient tools for predictive data analysis.
- Accessible to everybody and reusable in various contexts.
- Built on NumPy, SciPy, and matplotlib.
- Open source, commercially usable - BSD license.

## Installation

Scikit-Learn requires:

- Python ( $\geq 3.6$ )
- NumPy ( $\geq 1.13.3$ )

- SciPy ( $\geq 0.19.1$ ) It can be installed using pip:

```
pip install -U scikit-learn
```

## Quick Start

---

A simple example of using Scikit-Learn is fitting a linear regression model:

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
predictions = model.predict(X_test)
```

# Getting Started

---

## Loading Datasets

---

Scikit-Learn provides utilities for loading standard datasets for practicing machine learning techniques. These can be accessed via the `datasets` module.

## Data Preprocessing

---

Data preprocessing involves scaling, normalization, encoding, and handling missing values, which can be done using modules such as `preprocessing` and `impute`.

## Supervised Learning

---

### Classification

Scikit-Learn offers a variety of algorithms for classification tasks, from simple linear classifiers to complex ensemble methods.

### Regression

For regression tasks, Scikit-Learn provides several methods ranging from linear regression to more advanced regression techniques.

# Unsupervised Learning

---

## Clustering

Clustering algorithms like K-Means, hierarchical clustering, and DBSCAN are available for identifying groups in data.

## Dimensionality Reduction

Techniques like PCA, t-SNE, and LDA for dimensionality reduction are essential for analyzing high-dimensional data.

## Model Selection and Evaluation

---

### Cross-Validation

Cross-validation techniques help in assessing the performance of models reliably.

### Hyperparameter Tuning

Scikit-Learn's `GridSearchCV` and `RandomizedSearchCV` are powerful tools for hyperparameter tuning to optimize model performance.

### Model Evaluation Metrics

The library provides various metrics and scoring functions for evaluating the performance of classification, regression, and clustering models.

## Pipelines

---

### Building Pipelines

Pipelines help in chaining preprocessors and models, simplifying the workflow of model training.

### Grid Search Pipelines

Combining pipelines with grid search enables efficient hyperparameter tuning across the entire model building process.

## Working with Text Data

---

## Text Feature Extraction

Scikit-Learn offers utilities for extracting features from text data, such as `CountVectorizer` and `TfidfVectorizer`.

## Text Classification

Text data can be classified using Scikit-Learn's algorithms, once appropriate feature extraction is performed.

# API Reference

---

## BaseEstimator

---

The `BaseEstimator` class provides basic functionality for all estimators in scikit-learn. Key methods include `fit`, `predict`, and `score`.

## TransformerMixin

`TransformerMixin` is a mixin class for all transformers in scikit-learn. It provides a `fit_transform` method, which applies a fit on the data followed by a transform.

## Model Classes

---

### Linear Models

Linear models include classes like `LinearRegression`, `LogisticRegression`, and `Ridge`. These models are used for tasks in both regression and classification.

### Ensemble Models

Ensemble models, including `RandomForestClassifier` and `GradientBoostingRegressor`, combine the predictions of several base estimators to improve generalizability and robustness.

### Support Vector Machines

Support Vector Machines (SVMs) like `SVC` and `SVR` are powerful linear models used for classification and regression tasks.

### Nearest Neighbors

The Nearest Neighbors algorithm is implemented in classes like `KNeighborsClassifier` and `KNeighborsRegressor`, used for classification and regression.

## Preprocessing

---

### StandardScaler

`StandardScaler` standardizes features by removing the mean and scaling to unit variance.

### MinMaxScaler

`MinMaxScaler` scales features to a given range, typically between zero and one.

### OneHotEncoder

`OneHotEncoder` encodes categorical features as a one-hot numeric array.

## Metrics

---

### Classification Metrics

Includes metrics like accuracy, precision, recall, and F1-score, implemented in `sklearn.metrics`.

### Regression Metrics

Covers metrics for regression tasks, such as mean squared error, mean absolute error, and R-squared.

## Cross-Validation

---

### KFold

`KFold` splits a dataset into K consecutive folds, providing train/test indices to split data in train/test sets.

### StratifiedKFold

,

`StratifiedKFold` is a variation of k-fold which returns stratified folds: each set contains approximately the same percentage of samples of each target class as the complete set.

# Contributing

---

Contributing to open-source projects like Scikit-Learn is a rewarding way to learn, teach, and build experience in just about any skill you can imagine.

## How to Contribute

---

- **Contributing Code:** Steps to contribute code, including setting up a development environment, forking the repository, and creating a pull request.
- **Contributing Documentation:** Guidelines for writing and submitting documentation improvements.
- **Community Involvement:** Participating in discussions on mailing lists, submitting bug reports, and proposing new features.

## Code of Conduct

---

The project adheres to a code of conduct that should be read and followed by all contributors. This includes respecting others, being considerate, and fostering an inclusive environment.

## Development Guidelines

---

- **Coding Standards:** Adhering to coding standards and best practices for Python and specific to the project.
- **Testing:** Guidelines for writing tests to ensure code reliability and functionality.
- **Documentation:** Standards for documenting new code and features, ensuring clarity and accessibility for users.

## Reporting Issues

---

How to report issues effectively:

- **Using the Issue Tracker:** How to use the project's issue tracker to report bugs or suggest enhancements.
- **Providing Reproducible Examples:** Importance of providing a minimal, reproducible example when reporting a bug.
- **Issue Classification:** Understanding the classification of issues into bugs, feature requests, or enhancements.

# Examples

---

This section provides examples to illustrate how Scikit-Learn can be used in different scenarios, ranging from basic to advanced applications, and including real-world use cases.

## Basic Examples

---

Basic examples cover fundamental operations and simple use cases to help new users understand the basics of Scikit-Learn.

- **Example 1: Basic Data Preprocessing:** Demonstrating data scaling and normalization.
- **Example 2: Simple Linear Regression:** Building and evaluating a basic linear regression model.
- **Example 3: Basic Classification with k-NN:** Implementing a k-Nearest Neighbors classifier.

## Advanced Examples

---

Advanced examples delve into more complex scenarios and sophisticated techniques.

- **Example 1: Hyperparameter Tuning:** Using GridSearchCV for optimizing model parameters.
- **Example 2: Pipeline Creation:** Building a pipeline for data preprocessing and model training.
- **Example 3: Ensemble Methods:** Implementing a Random Forest classifier and understanding its advantages.

## Real-world Use Cases

---

Illustrating how Scikit-Learn can be applied to solve real-world problems.

- **Use Case 1: Image Recognition:** Applying a convolutional neural network for image classification.
- **Use Case 2: Natural Language Processing:** Implementing text classification using TF-IDF and logistic regression.
- **Use Case 3: Predictive Maintenance:** Using regression techniques for predicting equipment failures.

## Glossary

---

This glossary section provides definitions of key terms used in the Scikit-Learn documentation and user guide.

# Key Terms

---

1. Estimator
2. Transformer
3. Predictor
4. Model
5. Fit
6. Train/Test Split
7. Cross-Validation
8. Hyperparameter
9. Feature
10. Target
11. Classification
12. Regression
13. Clustering
14. Pipeline
15. Grid Search

# Definitions

















---

- **Estimator:** A class in Scikit-Learn that provides a `fit` method to learn from data.
- **Transformer:** A class that provides a `transform` method to modify or filter data.
- **Predictor:** A class that provides a `predict` method to make predictions based on learned parameters.
- **Model:** A statistical representation of a concept or process, used for making predictions or understanding data.
- **Fit:** The process of determining model parameters based on training data.
- **Train/Test Split:** The method of dividing a dataset into separate sets for training and testing a model.
- **Cross-Validation:** A technique for evaluating a model by partitioning the original sample into a training set to train the model, and a test set to evaluate it.
- **Hyperparameter:** A parameter of a learning algorithm (not of the model) which is set prior to the learning process.
- **Feature:** An individual measurable property or characteristic of a phenomenon being observed.
- **Target:** The output variable in a supervised learning algorithm.
- **Classification:** A type of supervised learning in which the goal is to predict categorical labels.



- **Regression:** A type of supervised learning in which the goal is to predict continuous values.
- **Clustering:** A type of unsupervised learning used to find groups in data.
- **Pipeline:** A tool for encapsulating multiple processing steps into a single unit.
- **Grid Search:** A method of tuning hyperparameters to find the best performance.

# Full Free Complete Artificial Intelligence Career Roadmap

Roadmap	Code	Documentation	Tutorial
1 TensorFlow Developers Roadmap	 Code <a href="#">TensorFlow Developers</a>	Docs <a href="#">TensorFlow</a>	 Tutorial <a href="#">TensorFlow</a>
2 PyTorch Developers Roadmap	 Code <a href="#">PyTorch Developers</a>	Docs <a href="#">PyTorch</a>	 Tutorial <a href="#">PyTorch</a>
3 Fundamentals of Computer Vision and Image Processing	 Code <a href="#">Computer Vision</a>	Docs <a href="#">OpenCV</a>	 Tutorial <a href="#">Computer Vision</a>
4 Statistics Roadmap for Data Science and Data Analysis	 Code <a href="#">Statistics</a>	Docs <a href="#">Statistics</a>	 Tutorial <a href="#">Statistics</a>
5 Becoming A Python Developer	 Code <a href="#">Python Developer</a>	Docs <a href="#">Python</a>	 Tutorial <a href="#">Python</a>
6 Machine Learning Engineer Roadmap	 Code <a href="#">Machine Learning Engineer</a>	Docs <a href="#">Machine Learning</a>	 Tutorial <a href="#">Machine Learning</a>
7 Become A Data Scientist	 Code <a href="#">Data Scientist</a>	Docs <a href="#">Data Science</a>	 Tutorial <a href="#">Data Science</a>
8 Deep Learning Engineer Roadmap	 Code <a href="#">Deep Learning Engineer</a>	Docs <a href="#">Deep Learning</a>	 Tutorial <a href="#">Deep Learning</a>